

Fact Data Modeling

What is a fact?

- Something that happened
 - User logins
 - Transactions that are made
 - Events (atomic pieces)

Fact Data 101:

- Fact data is 10-100x the volume of dim data
 - X actions per dimension
- Context for effective analysis is key
- Duplicates are more common (dedupe needed)
 - challenging!

Normalized vs. Denormalized Facts

- Normalized facts have no additional dim attributes (JOINS on ID)
- Denormalized facts store some dim data for faster analytics but higher storage
- The smaller the scale the better Normalization will be

Raw logs

- Used by Software Engineers
- Potential quality errors
- Short retention

Married
↔

Fact data

- Quality guarantees
 - Nice schema
 - Longer retention
- } Highly trusted

Modeling

- „Who“ fields are usually represented as ID
- „Where“ also IDs but is likely to bring in dims (especially high cardinality)
- „How“ similar to Where field
- „What“ fields are atomic events and fundament of fact data
- „When“ fields defines timestamp of the event and is also fundamentally part of fact data
 - ↳ Use UTC!

Key properties

- Data Quality guarantees - no duplicates, what and when not null
- Smaller than raw logs
- Columns should be easy to query and understand

Logging

- Only log what you really need
- Use shared schema in logs and pipeline
 - ↳ middle layer like Thrift

Options for high volume facts

Sampling

- metric driven use cases
- not when precision is needed

Bucketing

- can be bucketed by one important dim („Who“ IDs)
- Sorted-merge bucket (SMB) joins without shuffle

Retention

- Dim data as long as allowed/needed
- Fact data very costly so retention needs to be defined
- do not decide for „just in case“

Deduplication

- Define a time frame to "care" about duplicates
- Streaming allows to capture most of the duplicates efficiently
 - 15-60 minute windows are sweet spot
- Group by every hour to remove duplicates
 - Fold over join 2 hour windows
 - JOIN those into 4 hour windows
 - ... → Tree structure

Facts vs. Dimensions

- fact is event driven like user is active
- dims are state driven
- facts can be aggregated into dims
- CASE WHEN to bucketize aggregated facts is useful to reduce cardinality
 - ↳ 5-10 buckets sweep spot

Dimensions

- Usually show up in Group By
- Can be high and low cardinality
- Come from a snapshot of state

Facts

- Usually aggregated in analytics (SUM, AVG, COUNT)
- Higher volume than dimensions (1 user can make multiple actions)
- come from events and logs

Categorical Fact/Dimensions

- Class category derived from facts
- CASE WHEN logic and

Date List data structure (extremely efficient)

- Example cumulated schema

- user_id

- Date

- Dates-active

turn
→
info

- User_id (32)

- Date (2023-01-01) -1 day

- `dateList_int` (10000¹⁰⁰⁰⁰)

current day active

Data Shuffling

- Should be minimized to enhance parallelism
- Bottleneck for Big Data Processing
- Spark will use shuffle partitions to distribute data (default 200)

Efficiency

- Improve GroupBy by bucketing data on object storage (pre shuffle)
- Reduce data volume as much as possible
 - ↳ Start aggregating every action into daily/monthly/yearly data
 - ↳ Some analytics flexibility gets lost
 - ↳ Time horizons of aggregated data increases for analytics

Long-Array Metrics

- Super Powerful for huge history
- Monthly metrics with an array column for every day
- start-date/month-start as an index/offset
- Position in array determines the day of the month
- Similar to datelist for non binary data
- You need to pick snapshots in time

Parallelism

- Extremely parallel
 - SELECT, FROM, WHERE
 - Kinda parallel
 - GROUP BY, JOIN, HAVING → all rows of Group By needs to be on one machine
 - Not parallel
 - ORDER BY → all data on one machine (use only at the end after agg)
- Exception → window function with partition by